

AI Does Not Have an Intelligence Problem. It Has a Responsibility Problem.

Notes from an ongoing experiment in governing AI-assisted work

By Douglas P. Galullo

Most public argument about artificial intelligence begins with a fear: that machines will take work away from people. The fear is reasonable. Employment, identity, and dignity are bound up in the work people do, and increasingly capable AI systems raise legitimate questions about which roles will remain human.

But studying the actual process of working with these systems — not the technology in the abstract, the daily work — pushed me toward a different question. What if AI's growing participation in consequential human work is not something society can simply prevent? The pace of development, the investment behind it, and the plain usefulness of automation all suggest that AI will keep entering business, government, research, and professional life. Particular uses can be slowed, prohibited, or fenced off. Stopping the larger movement is much harder to imagine.

If that participation is close to inevitable, then protecting people cannot rest entirely on preserving every task in its current form. Protection must also come from governing how AI enters the work: what authority it receives, what evidence it uses, how its output is reviewed, and who remains responsible for the result.

I want to be precise about where that conclusion came from, because it did not come from philosophy or a position I went looking to confirm. It emerged slowly from studying a working process — observing it, changing it deliberately, and studying what the changes produced. The theory in this article is a residue of that iteration: what remained after the cycles were run.

Studying the Work, Not the Technology

My work with AI did not start as research. It started as ordinary tool use. I was using several AI systems to research, write, analyze, challenge assumptions, organize information, and develop software and operating procedures.

The lens I brought to that use was formed long before AI. I spent years inside operational environments — manufacturing, postal operations, business ownership, law enforcement — where the work itself trains certain habits: studying how a process actually flows rather than how it is described; deciding quickly with incomplete information; knowing who holds authority; preserving evidence; handling exceptions; watching what a decision does two and three steps downstream; and

answering personally for real-world outcomes. None of this makes anyone a better judge of algorithms. It does make process visible in a way that raw output is not. So when I began using AI systems seriously, I found myself watching the process around them at least as closely as the answers they produced.

Two things became clear quickly. The first was that “AI” is not one uniform capability. Different systems behaved differently: one was strong at structured reasoning but careless with the original intent of an assignment; another produced better prose but missed logical defects; a third was most useful when asked to attack an argument. The lesson was not that one system was better than another. It was that each was a tool with a particular range of appropriate uses and limits.

Nobody walks into a hardware store, buys the most powerful tool on the shelf, and expects it to build a house. Every tool in the bag is necessary; none substitutes for the others. Choosing the right tool is only the beginning: the builder still has to know what is being built, put each tool where it belongs, inspect the work, correct the defects, and answer for the finished structure. AI deserves the same discipline. Capability does not establish suitability. Suitability does not establish authority. Output does not establish truth. And the presence of a human somewhere in the workflow does not, by itself, establish oversight.

The second thing that became clear was subtler. Once several systems participated in the same body of work, combined capability rose — and so did the complexity around them. Information fragmented across tools. Context had to be rebuilt again and again. A conclusion would drift loose from the evidence that had originally supported it. Conflicting answers demanded interpretation. In time it became genuinely difficult to reconstruct which system had done what, why a conclusion was accepted, or where human judgment had entered.

The operational habit of following consequences shaped what I noticed next: the examination rarely stopped at the immediate output, because in consequential work that is never where the story ends. A plausible answer today becomes tomorrow’s assumption, and an accumulated assumption eventually becomes a procedure nobody remembers adopting. Authority granted for one task quietly transfers to the next. Evidence supporting a conclusion goes missing by the time the conclusion is acted on. Small departures from procedure drift into unexamined defaults. Consequential work cannot be evaluated solely by the quality of an immediate answer. It has to be evaluated by what the answer sets in motion — the second- and third-order effects, the downstream operational load, and the long-term question of who remains responsible.

Seen that way, the problem was no longer whether any individual answer was good or bad, but the process connecting the tools, the evidence, the decisions, the records, and the human authority over all of it.

Turning Friction into a Method

Studying that process required making it visible. A process can only be examined once its parts and their movement can be seen, and randomly comparing outputs from unrelated systems is not an experiment. So the first intervention was simply to define the connected process: work enters, assignments are created, evidence is reviewed, conflicts appear, decisions are made, actions follow, and records are preserved — or lost.

Once the process was visible, friction could be located, and a working method took shape: observe; identify friction; introduce a deliberate structural change; run the same kind of work again; evaluate what improved and what worsened; keep the useful change; correct the failure; observe again. The experiment was never a one-time comparison of AI systems. The experiment was the evolving process itself — studied, changed, studied again. The examination method deserves a fuller treatment than this article attempts; here it matters only as the route by which the responsibility thesis emerged.

Certain failures kept recurring, and each pointed to a missing function. Lost context pointed to persistent institutional memory. Repeated explanation pointed to reusable structured knowledge. Conflicting outputs pointed to deliberate challenge and an identified decision authority. Conclusions separated from their evidence pointed to preserving the link between claim and support. Ambiguous approvals pointed to distinguishing a suggestion from a proposal, a proposal from a decision, and a decision from an authorized action. Unbounded activity pointed to separating capability from permission. And nominal human involvement demonstrated that a person “in the loop” is not the same thing as a person responsible.

The response to those failures was not to hand any one system more control. It was to build structure around the work, drawing on an analogy to the way mature organizations separate responsibilities. Such institutions define authority, separate functions, preserve records, and create channels for review, dissent, escalation, and final approval.

AI-assisted work, it turned out, needed comparable discipline. Different systems began to receive bounded roles rather than open-ended authority. A coordinating function — something like a chief of staff — maintained continuity and routed assignments to the right specialist. A review function, resembling an independent challenge process, kept the first plausible answer from automatically becoming the accepted one. A recordkeeping function preserved decisions, approvals, changes, and unresolved conflicts, making the work reconstructable. An identified human remained responsible for consequential decisions and for accepting the results. None of this was conceived as a feature. Each element was a repair, introduced because the previous version of the workflow had failed in an identifiable way.

The structure also changed what human oversight meant. An approval is weak if the reviewer cannot see what was done, what evidence and limits were known, or what objections were raised. Clicking “approve” at the end of an opaque process may transfer blame; it does not demonstrate judgment. Meaningful responsibility requires a reconstructable record: what was reviewed, under what

authority, what alternatives were weighed, and what responsibility was accepted.

The plain lesson underneath all of it: intelligence does not govern itself.

Procedures That Learn

As the process evolved, standard operating procedures became central — but not in the usual sense of fixed instructions written once and followed indefinitely. That model is inadequate alongside systems whose capabilities change by the month.

In this study, an SOP became something else: a controlled, living record of the best-known way to perform a process at a particular moment in time. Friction is identified; a change is proposed and tested; if it improves the work, the SOP is revised and the previous version preserved. The process learns through experience. The SOP preserves the learning. Governance controls how that learning becomes the new standard. Without this, learning stays informal: people follow different versions of the process, systems receive inconsistent instructions, exceptions harden into habits, and corrected failures recur.

But adaptability cannot mean improvisation. A mature process has to distinguish among four things that look similar from the outside: an approved evolution, a temporary exception, the correction of a defective rule, and an unauthorized departure from procedure. A meaningful change should preserve what changed, why, on what evidence, who reviewed and authorized it, when it took effect, and which version it replaced. That history is not administrative clutter; it is evidence. A process that changes without recording why loses the ability to tell whether the new procedure solved the original problem or merely created a different one.

This is what adaptive governance means in practice. Rigid governance eventually becomes obsolete; ungoverned change eventually becomes arbitrary. The aim is neither. New conditions will keep arriving — new capabilities, new conflicts, failure modes that surface only with use — and the process should be able to recognize each one, contain it, decide whether it is an exception or a structural problem, and change the governing procedure when warranted. No consequential process becomes permanently perfect. The goal is a process that grows steadily more reliable, measurable, explainable, and resistant to known failures while remaining able to adapt to unknown ones. The principles stay stable while the procedures evolve: authority defined, evidence connected to claims, review meaningful, exceptions visible, responsibility identifiable.

Responsibility Infrastructure

Only after many of these cycles did the study produce a broader hypothesis. The central problem in consequential AI work is not whether a system can perform a task. It is whether the surrounding process can answer basic questions. Who authorized the work? What was the AI permitted to do? Which system performed which function? What evidence supported the conclusion? What

limitations were known? What alternatives or objections were considered, and where was the work challenged? What decision was ultimately made, who approved it, and who accepted responsibility for the outcome? Can the whole sequence be reconstructed later?

Ordinary activity logs cannot answer most of these, and the reason is structural. A log records events: a prompt was entered, an output produced, a file changed, a button clicked. Governance concerns warrants: whether the activity was authorized, whether the evidence was sufficient, whether required review occurred, whether the reviewer exercised judgment. A record of activity is not a record of authority.

Consequential AI work therefore needs more than an activity trail. It needs a responsibility structure, which can be compressed into a sentence: authority enters through a mandate; evidence and decisions persist through a governed record; responsibility exits through explicit acceptance.

This is not a proposal to treat an AI system as a legal or moral person. The point runs the other way: to prevent AI participation from dissolving the responsibility of the humans and institutions around it. If a system is given authority, that authority should come from somewhere identifiable. If it contributes to a consequential recommendation, its evidence and its limits should stay visible. If a human accepts the recommendation, that acceptance should represent an actual decision — not the final click in a process no one can reconstruct.

I have been testing this hypothesis inside an experimental governed operating environment called Doghouse, which matters to this argument not as a product but as the setting where the ideas meet real work. The study remains an ongoing, iterative experiment. It has not established that one structure fits every organization or industry. It has not solved AI accountability. What it has done is make the process visible enough to study — and that alone changes the quality of the questions. Instead of asking only whether an output looks correct, it can ask whether the work was authorized, evidenced, challenged, and accepted.

The obvious objection is bureaucracy. Governance can slow work down; badly designed, it produces paperwork without insight. The answer is proportion: a low-risk drafting task should not carry the same process as a decision affecting health, money, legal rights, or safety. Governance should rise with consequence; done well, it reduces friction rather than merely adding checkpoints. Persistent memory reduces repeated explanation, defined roles reduce duplicated effort, preserved evidence speeds review, and controlled procedures keep the same failure from recurring. The question is never whether governance adds steps. It is whether those steps cost less than the confusion, rework, and untraceable responsibility they replace.

A second objection cuts deeper: humans may not understand advanced systems well enough to review them meaningfully. That concern is legitimate — human approval is not valuable merely because it is human — but the answer cannot be ceremonial oversight. The process has to improve what the reviewer can actually see — the evidence, the disclosed limitations, the unresolved

conflicts, the path by which a recommendation formed. No structure guarantees understanding. A good one can at least prevent organizations from pretending that an opaque approval represents responsibility.

The race for capability will continue; models will become faster, more autonomous, more deeply embedded in consequential work. But greater capability will not answer the questions its use creates. Who gave the system authority? What evidence did it rely on? Who challenged its conclusion? Who accepted responsibility for what happened next? These are not questions of intelligence. They are questions of process, institutions, and responsibility.

The next significant layer of AI development, in other words, may not be another more capable model. It may be the responsibility infrastructure required to make AI-assisted work governable, defensible, traceable, and accountable — from the beginning of the work through every consequence that follows. That infrastructure has no settled form yet, and this study is one small, unfinished attempt to find it. But its working hypothesis has held up through every iteration so far: AI will not become trustworthy merely by becoming more intelligent. Consequential AI work will become governable only when responsibility becomes part of the process itself.